Day 3

Model에게 원하는 것?

- Regression -> nn.MSELoss()
- Binary classification -> nn.BCELoss()
- Multi class classification -> nn.CrossEntropyLoss()

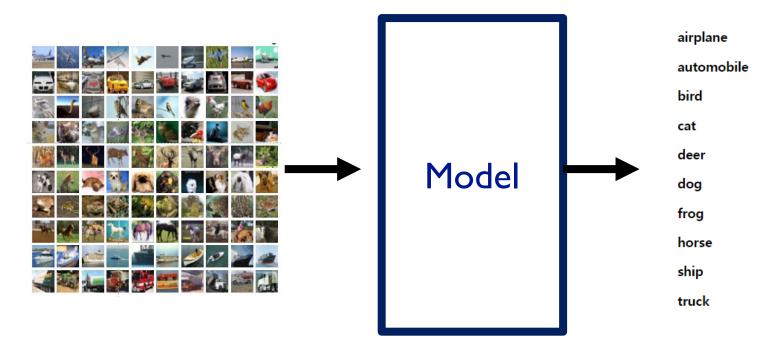
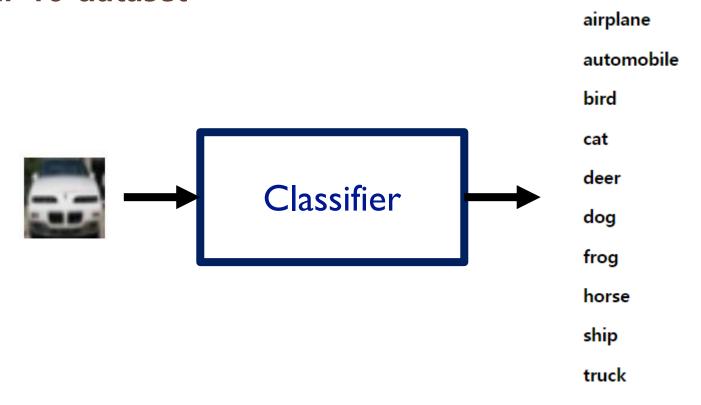


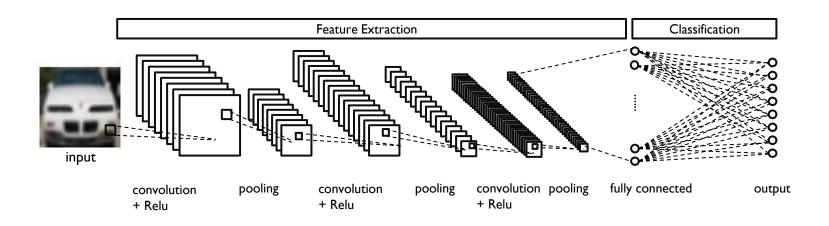
Image Classification

Cifar 10 dataset



Coding 전에 생각해 볼 것

- ▶ 입력 데이터 : image(C , H , W)
- > 출력 데이터 : class number
- Optimizer: ?
- Loss function: cross entropy



1. Dataset

▶ 입, 출력 데이터

```
transform_CIFAR10 = transforms.Compose([
   transforms.ToTensor()
train_loader = torch.utils.data.DataLoader(
   datasets.CIFAR10(
       root = './data_CIFAR10',
       train = True,
       download = True,
       transform = transform_CIFAR10),
   batch_size=BATCH_SIZE,
   shuffle=True
test_loader = torch.utils.data.DataLoader(
   datasets.CIFAR10(
       root = './data_CIFAR10',
       train = False.
       download = True,
       transform = transform_CIFAR10),
   batch_size=BATCH_SIZE,
   shuffle=True
```

2. Model

► CNN + Linear

```
class CNN(nn Module):
   def __init__(self):
       super(CNN, self).__init__()
       self.conv1 = nn.Conv2d(3, 8, kernel_size=3)
       self.conv2 = nn.Conv2d(8, 16, kernel_size=3)
       self.conv3 = nn.Conv2d(16, 24, kernel size=3)
       #self.conv3 drop = nn.Dropout2d()
       self.fc1 = nn.Linear(24 * 2 * 2, 128)
       self.fc2 = nn.Linear(128, 10)
   def forward(self, x):
       x = F.relu(F.max.pool2d(self.conv1(x), 2))
       x = F.relu(F.max_pool2d(self.conv2(x), 2))
       x = F.relu(F.max_pool2d(F.dropout(self.conv3(x)), 2))
       x = x.view(-1, 24 * 2 * 2)
       x = F.relu(self.fcl(x))
       x = F.dropout(x, training=self.training)
       x = self.fc2(x)
       return x
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 8, 30, 30]	224
Conv2d-2	[-1, 16, 13, 13]	1,168
Conv2d-3	[-1, 24, 4, 4]	3,480
Dropout2d-4	[-1, 24, 4, 4]	0
Linear-5	[-1, 128]	12,416
Linear-6	[-1, 10]	1,290

Total params: 18,578 Trainable params: 18,578 Non-trainable params: 0

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 8, 30, 30]	224
Conv2d-2	[-1, 16, 13, 13]	1,168
Conv2d-3	[-1, 24, 4, 4]	3,480
Linear-4	[-1, 128]	12,416
Linear-5	[-1, 10]	1,290

Total params: 18,578 Trainable params: 18,578 Non-trainable params: 0

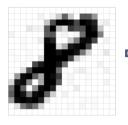


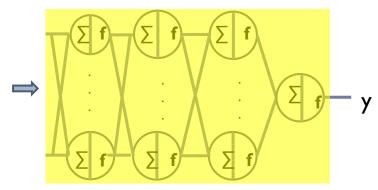
nn.Linear

```
RECAP
```

```
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.fc1 = nn.Linear(784, 256)
        self.fc2 = nn.Linear(256, 128)
        self.fc3 = nn.Linear(128, 10)

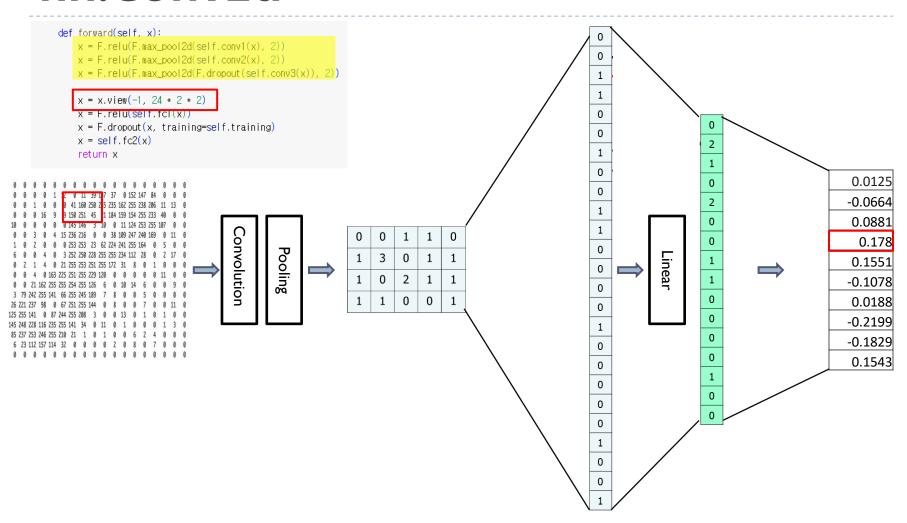
def forward(self, x):
        x = x.view(-1, 784)
        x = torch.sigmoid(self.fc1(x))
        x = self.fc3(x)
    return x
```



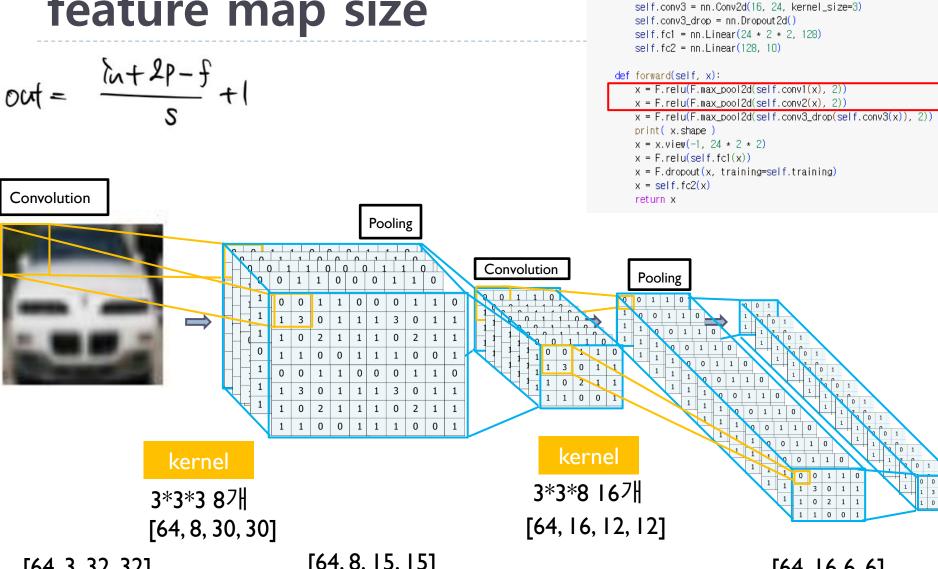




nn.Conv2d



feature map size



class CNN(nn Module): def __init__(self):

super(CNN, self).__init__()

self.conv1 = nn.Conv2d(3, 8, kernel_size=3) self.conv2 = nn.Conv2d(8, 16, kernel_size=3)

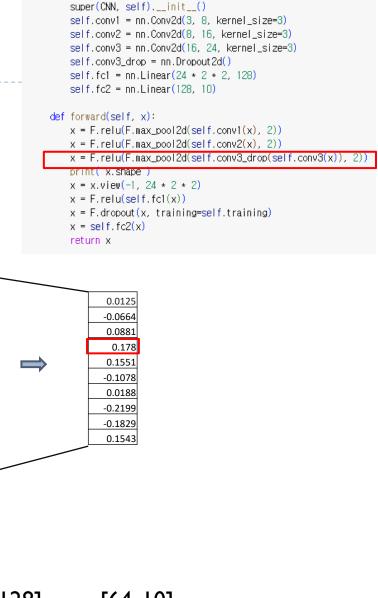
[64, 3, 32, 32]

[64, 8, 15, 15]

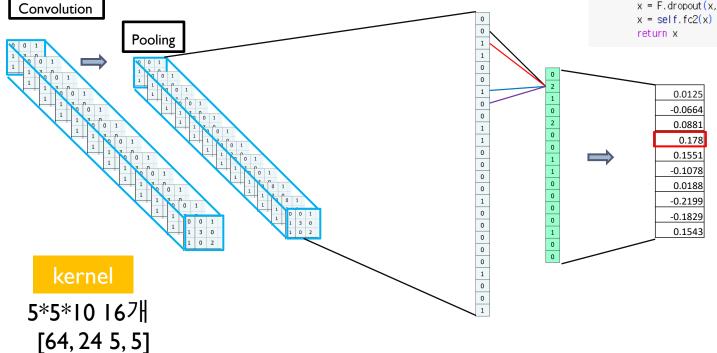
[64, 16 6, 6]



feature map size



class CNN(nn Module): def __init__(self):



[64, 24, 2, 2]

[64, 24*2*2] [64, 128]

[64, 10]



MaxPool2d

For each block, or "pool", the operation simply involves computing the [latex]max[/latex] value, like this:

4	3	1	5
1	3	4	8
4	5	4	3
6	5	9	4

$$Max([4, 3, 1, 3]) = 4$$

Doing so for each pool, we get a nicely downsampled outcome, greatly benefiting the spatial hierarchy we need:

4	3	1	5
1	3	4	8
4	5	4	3
6	5	9	4





Docs > torch.nn > MaxPool2d

>_

MAXPOOL2D

CLASS torch.nn.MaxPool2d(kernel_size, stride=None, padding=0, dilation=1, return_indices=False, ceil_mode=False) [SOURCE]

Applies a 2D max pooling over an input signal composed of several input planes.

In the simplest case, the output value of the layer with input size (N,C,H,W), output (N,C,H_{out},W_{out}) and kexnel_size (kH,kW) can be precisely described as:

$$out(N_i, C_j, h, w) = \max_{m=0,\dots,kH-1} \max_{n=0,\dots,kW-1} \max_{\text{input}(N_i, C_j, \text{stride}[0] \times h + m, \text{stride}[1] \times w + n)$$

If padding is non-zero, then the input is implicitly padded with negative infinity on both sides for padding number of points. dilation controls the spacing between the kernel points. It is harder to describe, but this link has a nice visualization of what dilation does.

https://github.com/christianversloot/machine-learning-articles/blob/main/what-are-max-pooling-average-pooling-global-max-pooling-and-global-average-pooling.md

AvgPool2d

Average Pooling

Another type of pooling layers is the Average Pooling layer. Here, rather than a max value, the avg for each block is computed:

4	3	1	5
1	3	4	8
4	5	4	3
6	5	9	4

$$Avg([4, 3, 1, 3]) = 2.75$$

As you can see, the output is also different - and less extreme compared to Max Pooling:





2.8	4.5
5.3	5.0

Docs > torch.nn > AvgPool2d

AVGPOOL 2D

Applies a 2D average pooling over an input signal composed of several input planes.

In the simplest case, the output value of the layer with input size (N,C,H,W), output (N,C,H_{out},W_{out}) and kernel_size (kH,kW) can be precisely described as:

$$out(N_i,C_j,h,w) = rac{1}{kH*kW}\sum_{m=0}^{kH-1}\sum_{n=0}^{kW-1}input(N_i,C_j,stride[0] imes h+m,stride[1] imes w+n)$$

If padding is non-zero, then the input is implicitly zero-padded on both sides for padding number of points.



>_

Global Max Pooling

Global Max Pooling

Another type of pooling layer is the Global Max Pooling layer. Here, we set the pool size equal to the input size, so that the max of the entire input is computed as the output value (Dernoncourt, 2017):

4	3	1	5
1	3	4	8
4	5	4	3
6	5	9	4

$$[4, 3, 1, 5]$$

$$[1, 3, 4, 8]$$

$$[4, 5, 4, 3]$$

$$[6, 5, 9, 4]$$

Or, visualizing it differently:

4	3	1	5
1	3	4	8
4	5	4	3
6	5	9	4



ADAPTIVEMAXPOOL2D

CLASS torch.nn.AdaptiveMaxPool2d(output_size, return_indices=False) [SOURCE]

Applies a 2D adaptive max pooling over an input signal composed of several input planes.

The output is of size $H_{out} imes W_{out}$, for any input size. The number of output features is equal to the number of input planes.

Parameters

- output_size the target output size of the image of the form $H_{out} \times W_{out}$. Can be a tuple (H_{out}, W_{out}) or a single H_{out} for a square image $H_{out} \times H_{out}$. H_{out} and W_{out} can be either a int, or None which means the size will be the same as that of the input.
- return_indices if True, will return the indices along with the outputs. Useful to pass to nn.MaxUnpool2d.
 Default: False

Shape:

- Input: (N, C, H_{in}, W_{in}) or (C, H_{in}, W_{in}) .
- Output: (N, C, H_{out}, W_{out}) or (C, H_{out}, W_{out}) , where $(H_{out}, W_{out}) = \text{output_size}$.

Examples

```
>>> # target output size of 5x7
>>> m = nn.AdaptiveMaxPool2d((5,7))
>>> input = torch.randn(1, 64, 8, 9)
>>> output = m(input)
>>> # target output size of 7x7 (square)
>>> m = nn.AdaptiveMaxPool2d(7)
>>> input = torch.randn(1, 64, 10, 9)
>>> output = m(input)
>>> # target output size of 10x7
>>> m = nn.AdaptiveMaxPool2d((None, 7))
>>> input = torch.randn(1, 64, 10, 9)
>>> output = m(input)
```

Global Avg Pooling

Global Average Pooling

When applying Global Average Pooling, the pool size is still set to the size of the layer input, but rather than the maximum, the average of the pool is taken:

4	3	1	5
1	3	4	8
4	5	4	3
6	5	9	4

Or, once again when visualized differently:

4	3	1	5
1	3	4	8
4	5	4	3
6	5	9	4





CLASS torch.nn.AdaptiveAvgPool2d(output_size) [SOURCE]

Applies a 2D adaptive average pooling over an input signal composed of several input planes.

The output is of size H x W, for any input size. The number of output features is equal to the number of input planes.

Parameters

output_size – the target output size of the image of the form H x W. Can be a tuple (H, W) or a single H for a square image H x H. H and W can be either a int, or None which means the size will be the same as that of the input.

Shape:

- Input: (N, C, H_{in}, W_{in}) or (C, H_{in}, W_{in}) .
- Output: (N,C,S_0,S_1) or (C,S_0,S_1) , where $S={
 m output_size}$.

Examples

```
>>> # target output size of 5x7
>>> m = nn.AdaptiveAvgPool2d((5,7))
>>> input = torch.randn(1, 64, 8, 9)
>>> output = m(input)
>>> # target output size of 7x7 (square)
>>> m = nn.AdaptiveAvgPool2d(7)
>>> input = torch.randn(1, 64, 10, 9)
>>> output = m(input)
>>> # target output size of 10x7
>>> m = nn.AdaptiveAvgPool2d((None, 7))
>>> input = torch.randn(1, 64, 10, 9)
>>> output = m(input)
```

ResNet



self.avgpool = nn.AdaptiveAvgPool2d((1, 1)

layer name	output size	18-layer	34-layer	50-layer
conv1	112×112			7×7, 64, stride
				3×3 max pool, stri
conv2_x	56×56	$\left[\begin{array}{c} 3\times3,64\\ 3\times3,64 \end{array}\right]\times2$	$\left[\begin{array}{c} 3\times3,64\\ 3\times3,64 \end{array}\right]\times3$	\[\begin{array}{c} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{array} \times 3
conv3_x	28×28	$\left[\begin{array}{c} 3\times3, 128\\ 3\times3, 128 \end{array}\right] \times 2$	$\left[\begin{array}{c} 3\times3,128\\ 3\times3,128 \end{array}\right]\times4$	\[\begin{array}{c} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{array} \times 4
conv4_x	14×14	$\begin{bmatrix} 3\times3, 256 \\ 3\times3, 256 \end{bmatrix} \times 2$	$\left[\begin{array}{c} 3\times3,256\\ 3\times3,256 \end{array}\right]\times6$	\[\begin{array}{c} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{array} \times 6 \]
conv5_x	7×7	$\left[\begin{array}{c} 3\times3,512\\ 3\times3,512 \end{array}\right]\times2$	$\left[\begin{array}{c}3\times3,512\\3\times3,512\end{array}\right]\times3$	\[\begin{array}{c} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{array} \times 3
	1×1		av	erage pool, 1000-d fc
FLO	OPs	1.8×10 ⁹	3.6×10 ⁹	3.8×10 ⁹

```
class ResNet(nn.Module):
    def __init__(self, block, num_blocks, num_classes=10):
        super(ResNet, self).__init__()
        self.in planes = 64
        self.conv1 = nn.Conv2d(3, 64, kernel_size=3,
                               stride=1, padding=1, bias=False)
        self.bn1 = nn.BatchNorm2d(64)
        self.layer1 = self. make layer(block, 64, num blocks[0], stride=1)
        self.layer2 = self. make layer(block, 128, num blocks[1], stride=2)
        self.layer3 = self._make_layer(block, 256, num_blocks[2], stride=2)
        self.layer4 = self._make_layer(block, 512, num_blocks[3], stride=2)
        self.linear = nn.Linear(512*block.expansion, num classes)
    def _make_layer(self, block, planes, num_blocks, stride):
        strides = [stride] + [1]*(num_blocks-1)
       layers = []
        for stride in strides:
            layers.append(block(self.in_planes, planes, stride))
            self.in_planes = planes * block.expansion
        return nn.Sequential(*layers)
    def forward(self, x):
        out = F.relu(self.bn1(self.conv1(x)))
        out = self.layer1(out)
        out = self.layer2(out)
        out = self.layer3(out)
        out = self.layer4(out)
        out = F.avg_pool2d(out, 4)
        out = out.view(out.size(0), -1)
        out = self.linear(out)
        return out
```



Question and Answer